

15-721 DATABASE SYSTEMS

Lecture #08 – Latch-free OLTP Indexes (Part II)

@Andy_Pavlo // Carnegie Mellon University // Spring 2017

TODAY'S AGENDA

Bw-Tree Index

ART Index

Profiling in Peloton



OBSERVATION

We cannot have reverse pointers in a latch-free concurrent Skip List because CaS can only update a single address at a time.

BW-TREE

Latch-free B+Tree index

→ Threads never need to set latches or block.

Key Idea #1: Deltas

→ No updates in place

→ Reduces cache invalidation.

Key Idea #2: Mapping Table

→ Allows for CAS of physical locations of pages.




THE BW-TREE: A B-TREE FOR NEW HARDWARE
ICDE 2013

BW-TREE: MAPPING TABLE

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	
104	●

Index Page



*Logical
Pointer* 


*Physical
Pointer* 

BW-TREE: MAPPING TABLE

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	
104	●

Index Page



*Logical
Pointer* 

*Physical
Pointer* 

BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

Page 102

*Logical
Pointer* 

*Physical
Pointer* 

Each update to a page produces a new delta.

BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

▲ Insert 50

Page 102

Logical
Pointer →

Physical
Pointer →

Each update to a page produces a new delta.

BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

▲ Insert 50

Page 102

Each update to a page produces a new delta.

Delta physically points to base page.

Logical Pointer ---→

Physical Pointer →

BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	

▲ Insert 50

Page 102

Logical
Pointer 

Physical
Pointer 

Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CAS.

BW-TREE: DELTA UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

▲ Insert 50

Page 102

Logical
Pointer ---→

Physical
Pointer →

Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CAS.

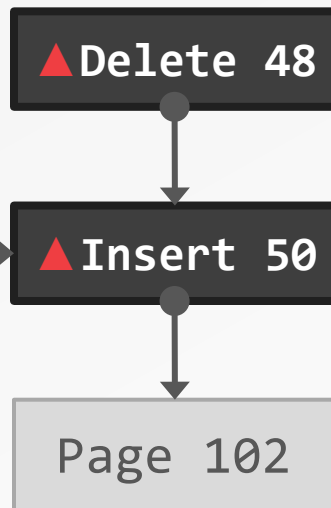
BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer —————>



Each update to a page produces a new delta.

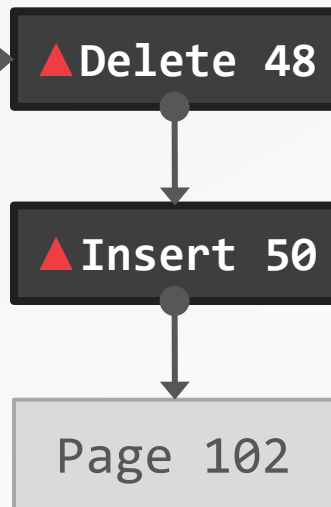
Delta physically points to base page.

Install delta address in physical address slot of mapping table using CAS.

BW-TREE: DELTA UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Each update to a page produces a new delta.

Delta physically points to base page.

Install delta address in physical address slot of mapping table using CAS.

Logical Pointer ----->


Physical Pointer ————>

BW-TREE: SEARCH

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* ----->
*Physical
Pointer* —————>

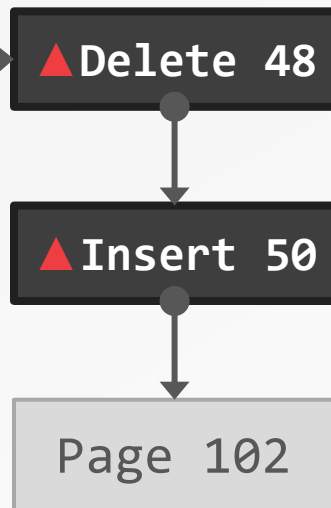


Traverse tree like a regular B+tree.

BW-TREE: SEARCH


Mapping Table


PID	Addr
101	
102	
103	
104	



Traverse tree like a regular B+tree.

If mapping table points to delta chain, stop at first occurrence of search key.

Logical
Pointer 

Physical
Pointer 

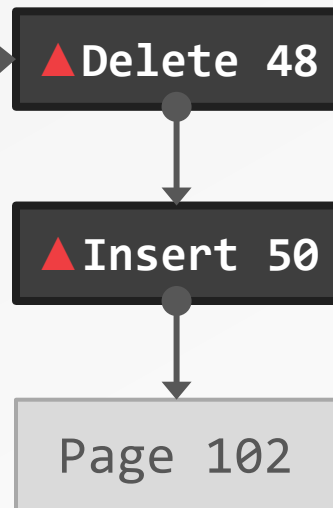
BW-TREE: SEARCH

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>



Traverse tree like a regular B+tree.

If mapping table points to delta chain, stop at first occurrence of search key.

Otherwise, perform binary search on base page.

BW-TREE: CONTENTION UPDATES

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

▲ Insert 50

Page 102

Threads may try to install updates to same state of the page.


Logical
Pointer ---→

Physical
Pointer →

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	●
103	
104	



Threads may try to install updates to same state of the page.


Logical
Pointer ---→

Physical
Pointer →

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same state of the page.

Winner succeeds, any losers must retry or abort

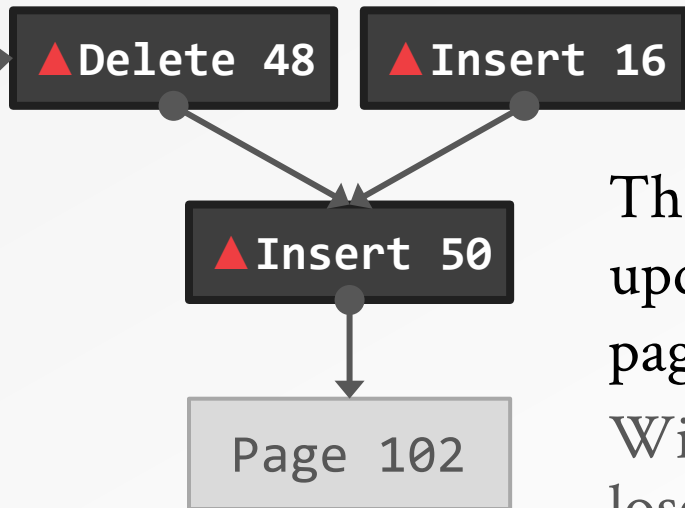
Logical
Pointer ---→

Physical
Pointer →

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same state of the page.

Winner succeeds, any losers must retry or abort


Logical Pointer ----->

Physical Pointer ————>

BW-TREE: CONTENTION UPDATES

Mapping Table

PID	Addr
101	
102	
103	
104	



Threads may try to install updates to same state of the page.

Winner succeeds, any losers must retry or abort

Logical Pointer ----->

Physical Pointer ————>

BW-TREE: DELTA TYPES

Record Update Deltas

→ Insert/Delete/Update of record on a page

Structure Modification Deltas

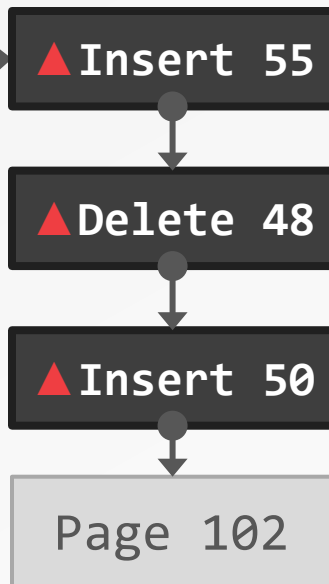
→ Split/Merge information




BW-TREE: CONSOLIDATION


Mapping Table

PID	Addr
101	
102	
103	
104	



Consolidate updates by creating new page with deltas applied.

Logical
Pointer 

Physical
Pointer 

BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102

Consolidate updates by creating new page with deltas applied.

Logical
Pointer ----->

Physical
Pointer ————>

New 102

BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102

Consolidate updates by creating new page with deltas applied.

Logical Pointer ----->

Physical Pointer ————>

New 102

▲ Insert 50

BW-TREE: CONSOLIDATION

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102

Consolidate updates by creating new page with deltas applied.
CAS-ing the mapping table address ensures no deltas are missed.

Logical Pointer ----->

Physical Pointer ————>

New 102

BW-TREE: CONSOLIDATION


Mapping Table

PID	Addr
101	
102	
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>

New 102



Consolidate updates by creating new page with deltas applied.
CAS-ing the mapping table address ensures no deltas are missed.

BW-TREE: CONSOLIDATION

Mapping Table

PID	Addr
101	
102	
103	
104	

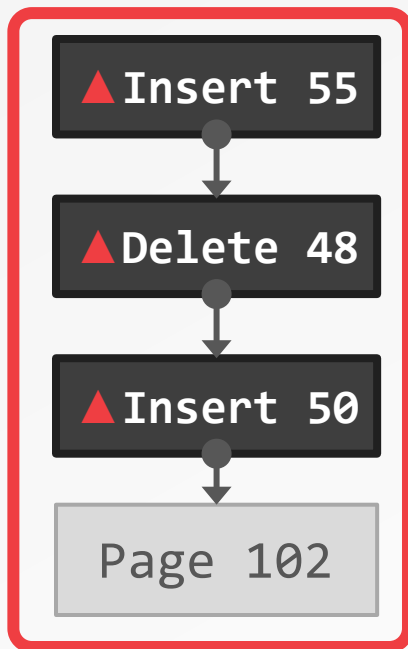
Logical
Pointer



Physical
Pointer



New 102



Consolidate updates by creating new page with deltas applied.

CAS-ing the mapping table address ensures no deltas are missed.

Old page + deltas are marked as garbage.

BW-TREE: GARBAGE COLLECTION

Operations are tagged with an **epoch**

- Each epoch tracks the threads that are part of it and the objects that can be reclaimed.
- Thread joins an epoch prior to each operation and post objects that can be reclaimed for the current epoch (not necessarily the one it joined)

Garbage for an epoch reclaimed only when all threads have exited the epoch.

BW-TREE: GARBAGE COLLECTION

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

*Logical
Pointer* ----->


*Physical
Pointer* —————>

▲ Insert 55

▲ Delete 48


▲ Insert 50

Page 102



New 102

Epoch Table



BW-TREE: GARBAGE COLLECTION

Mapping Table


<i>PID</i>	<i>Addr</i>
101	
102	
103	
104	

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102



CPU1

*Logical
Pointer* ---→


*Physical
Pointer* →

New 102



CPU2

Epoch Table



CPU1



CPU2

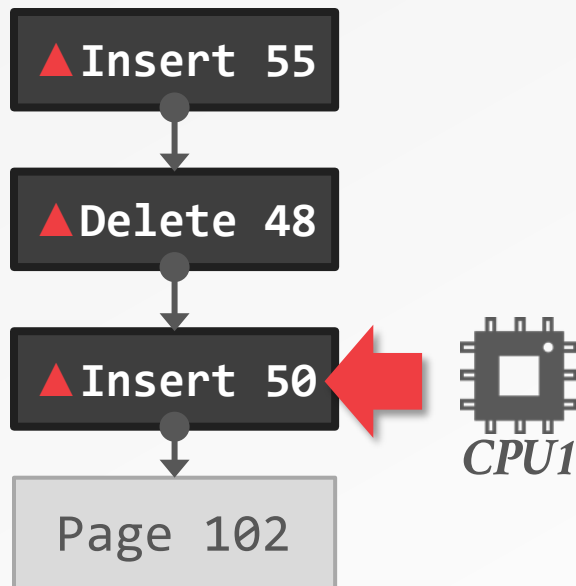
BW-TREE: GARBAGE COLLECTION

Mapping Table


PID	Addr
101	
102	●
103	
104	

Logical
Pointer ---→

Physical
Pointer →



Epoch Table



BW-TREE: GARBAGE COLLECTION

Mapping Table

PID	Addr
101	
102	●
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>


New 102

▲ Insert 55

▲ Delete 48


▲ Insert 50

Page 102



CPU1

Epoch Table



CPU1

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102


BW-TREE: GARBAGE COLLECTION

Mapping Table

<i>PID</i>	<i>Addr</i>
101	
102	●
103	
104	

*Logical
Pointer* ----->

*Physical
Pointer* ————>



New 102

Epoch Table



BW-TREE: GARBAGE COLLECTION

Mapping Table

PID	Addr
101	
102	●
103	
104	

Logical
Pointer ----->

Physical
Pointer ————>

New 102

Epoch Table

▲ Insert 55

▲ Delete 48

▲ Insert 50

Page 102

BW-TREE: STRUCTURE MODIFICATIONS

Split Delta Record

- Mark that a subset of the base page's key range is now located at another page.
- Use a logical pointer to the new page.

Separator Delta Record

- Provide a shortcut in the modified page's parent on what ranges to find the new page.


BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	

*Logical
Pointer* ---→


*Physical
Pointer* —→





BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	

Logical
Pointer 

Physical
Pointer 




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	

*Logical
Pointer* ---→


*Physical
Pointer* —→





BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	●

*Logical
Pointer* 

*Physical
Pointer* 




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---→

Physical
Pointer —→




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	●

Logical Pointer ---→

Physical Pointer →




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

<i>PID</i>	<i>Addr</i>
101	●
102	●
103	●
104	●
105	●

Logical Pointer ---→

Physical Pointer →




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---

Physical
Pointer —




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

*Logical
Pointer* ---→

*Physical
Pointer* —→




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---→

Physical
Pointer —→




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---

Physical
Pointer —




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical
Pointer ---→

Physical
Pointer —→




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

*Logical
Pointer* ---→

*Physical
Pointer* →




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical Pointer ---

Physical Pointer —




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●

Logical Pointer ---

Physical Pointer —




BW-TREE: STRUCTURE MODIFICATIONS

Mapping Table

PID	Addr
101	●
102	●
103	●
104	●
105	●


*Logical
Pointer* ---→

*Physical
Pointer* →



BW-TREE: PERFORMANCE

Processor: 1 socket, 4 cores w/ 2xHT



ADAPATIVE RADIX TREE (ART)

Uses digital representation of keys to examine prefixes one-by-one instead of comparing entire key.

Radix trees properties:


- The height of the tree depends on the length of keys.
- Does not require rebalancing
- The path to a leaf node represents the key of the leaf
- Keys are stored implicitly and can be reconstructed from paths.



THE ADAPTIVE RADIX TREE: ARTFUL
INDEXING FOR MAIN-MEMORY DATABASES
ICDE 2013

TRIE VS. RADIX TREE


Trie



Keys: **HELLO**, HAT, HAVE

TRIE VS. RADIX TREE


Trie



Keys: **HELLO**, HAT, HAVE

TRIE VS. RADIX TREE


Trie




Keys: HELLO, **HAT, HAVE**

TRIE VS. RADIX TREE

Trie




Radix Tree




Keys: HELLO, HAT, HAVE

TRIE VS. RADIX TREE

Trie




Radix Tree

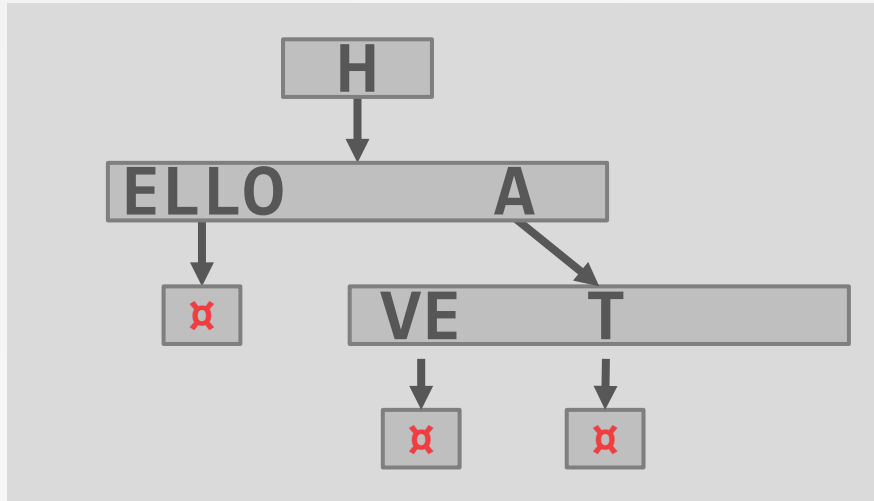


Keys: **HELLO**, HAT, HAVE

ART INDEX: MODIFICATIONS



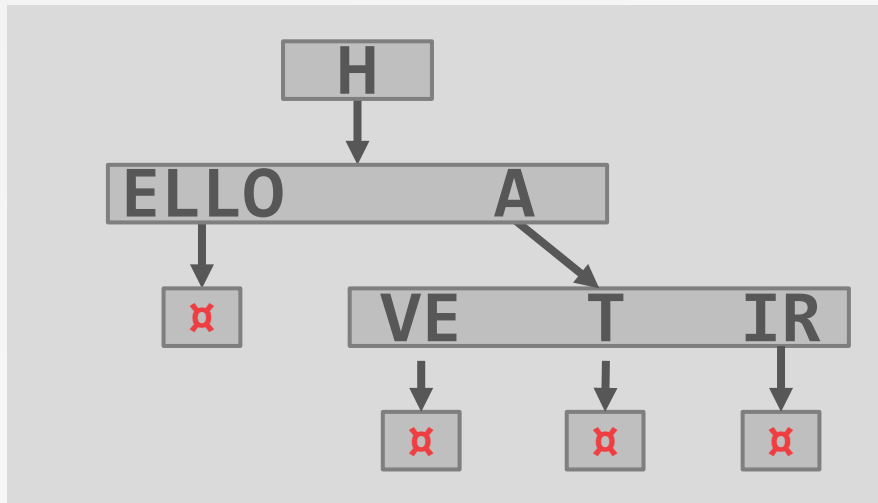
ART INDEX: MODIFICATIONS



Operation: Insert HAIR

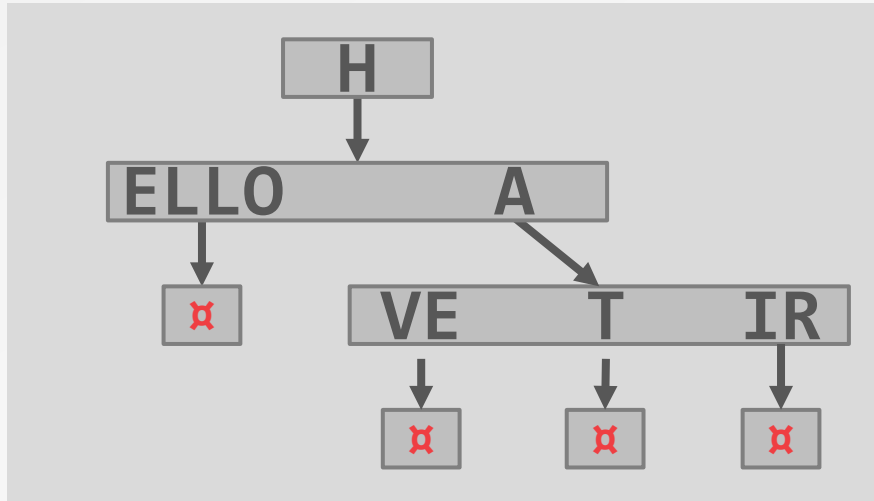


ART INDEX: MODIFICATIONS



Operation: Insert **HAIR**

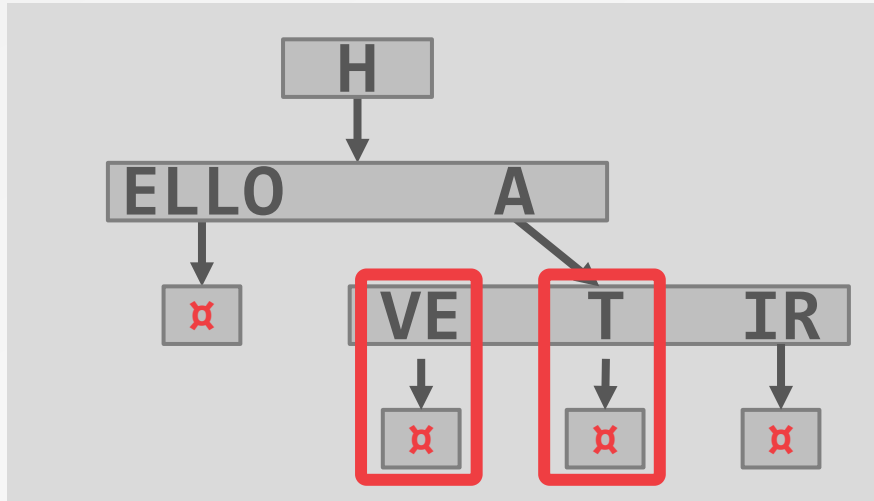
ART INDEX: MODIFICATIONS



Operation: Insert **HAIR**

Operation: Delete **HAT**, **HAVE**

ART INDEX: MODIFICATIONS




Operation: Insert **HAIR**

Operation: Delete **HAT**, **HAVE**



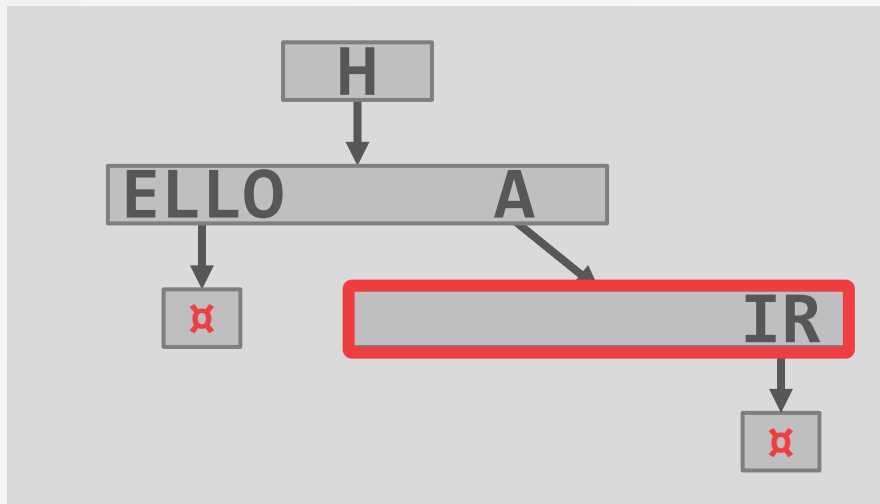
ART INDEX: MODIFICATIONS



Operation: Insert **HAIR**

Operation: Delete **HAT**, **HAVE**


ART INDEX: MODIFICATIONS



Operation: Insert **HAIR**

Operation: Delete **HAT**, **HAVE**

ART INDEX: MODIFICATIONS



Operation: Insert **HAIR**

Operation: Delete **HAT**, **HAVE**



ART INDEX: BINARY COMPARABLE KEYS

Not all attribute types can be decomposed into binary comparable digits for a radix tree.


- **Unsigned Integers:** Byte order must be flipped for little endian machines.
- **Signed Integers:** Flip two's-complement so that negative numbers are smaller than positive.
- **Floats:** Classify into group (neg vs. pos, normalized vs. denormalized), then store as unsigned integer.
- **Compound:** Transform each attribute separately.

ART INDEX: BINARY COMPARABLE KEYS

Int Key: 168496141



Hex Key: 0A 0B 0C 0D

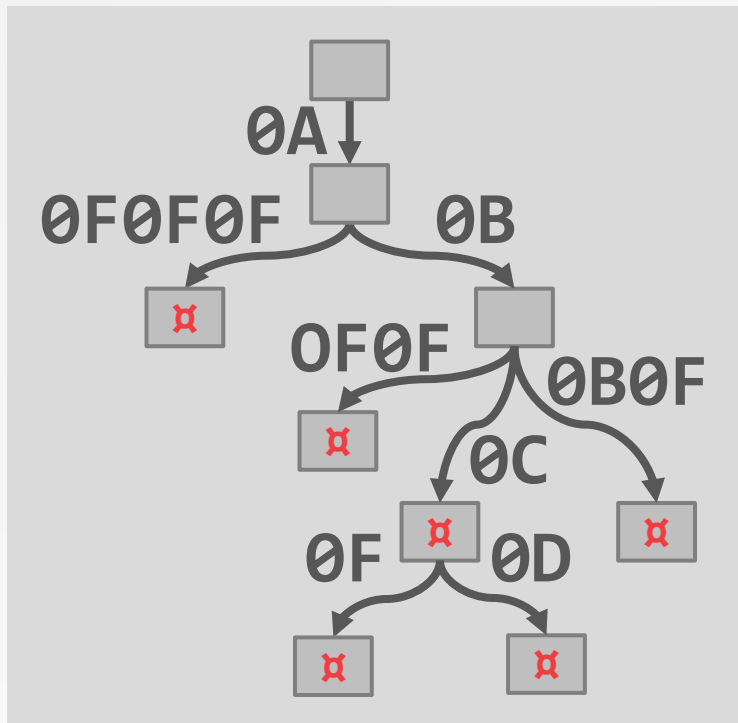


Big Endian



Little Endian


ART INDEX: BINARY COMPARABLE KEYS



Int Key: 168496141



Hex Key: 0A 0B 0C 0D

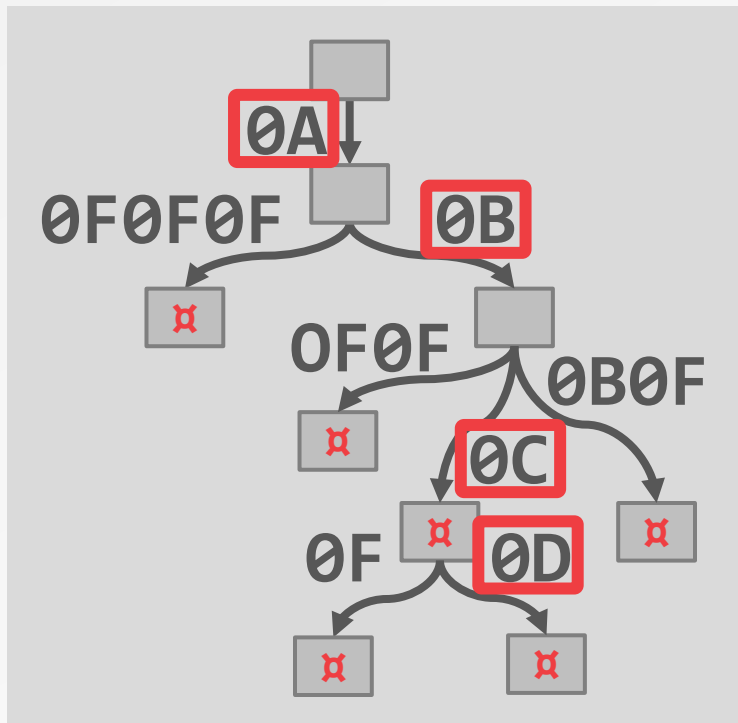


Big Endian



Little Endian


ART INDEX: BINARY COMPARABLE KEYS



Int Key: 168496141



Hex Key: 0A 0B 0C 0D




Big Endian



Little Endian

BINARY COMPRABLE KEYS

Peloton w/ Bw-Tree Index
Data Set: 10m keys (three 64-bit ints)



CONCURRENT ART INDEX

HyPer's ART is not latch-free.

Optimistic crabbing scheme where writers are not blocked on readers.


- Writers increment counter when they acquire latch.
- Readers can proceed if a node's latch is available.
- It then checks whether the latch's counter has changed from when it checked the latch.



THE ART OF PRACTICAL SYNCHRONIZATION
DaMoN 2016

SINGLE-THREADED PERFORMANCE

Data Set: 30m Random 64-bit Integers



PARTING THOUGHTS

Bw-Tree is probably the most dank latch-free index in recent years.

ART has amazing performance. Need to understand it better.





ANDY'S
**TIPS FOR
PROFILING**



MOTIVATION

Consider a program with functions **foo** and **bar**.

How can we speed it up with only a debugger ?

- Randomly pause it during execution
- Collect the function call stack



RANDOM PAUSE METHOD

Consider this scenario

- Collected 10 call stack samples
- Say 6 out of the 10 samples were in **foo**

What percentage of time was spent in **foo**?

- Roughly 60% of the time was spent in **foo**
- Accuracy increases with # of samples



AMDAHL'S LAW

Say we optimized **foo** to run 2 times faster

What's the expected overall speedup ?

→ 60% of time spent in **foo** drops in half

→ 40% of time spent in **bar** unaffected

→ **p** = percentage of time spent in optimized task

→ **s** = speed up for the optimized task

→ Overall speedup = $\frac{1}{\frac{p}{s} + (1-p)}$ = 1.4 times faster

AMDAHL'S LAW

1 0.6 2 +0.4 1 1 0.6 2 +0.4 0.6 2 0.6 0.6 2 2 0.6
2 +0.4 1 0.6 2 +0.4 = 1.4 times faster

1 0.6 2 +0.4 1 1 0.6 2 +0.4 0.6 2 0.6 0.6 2 2 0.6
2 +0.4 1 0.6 2 +0.4 = 1.4 times faster

$$\frac{1}{2 + 0.4} + (1 - \frac{1}{2 + 0.4}) \cdot 1 = 1.4 \text{ times faster}$$

Say we optimized **foo** to run 2 times faster

What's the expected overall speedup ?

- 60% of time spent in **foo** drops in half
- 40% of time spent in **bar** unaffected

PROFILING TOOLS FOR REAL

Choice #1: Valgrind

- Heavyweight instrumentation framework with a lot of tools
- Sophisticated visualization tools

Choice #2: Perf

- Lightweight tool that can record different kinds of events
- Console-oriented visualization tools

CHOICE #1: VALGRIND

Instrumentation framework for building dynamic analysis tools

- **memcheck**: a memory error detector
- **callgrind**: a call-graph generating profiler

Using **callgrind** to profile the index test and Peloton in general:

```
$ valgrind --tool=callgrind --trace-children=yes  
./tests/skiplist_index_test  
  
$ valgrind --tool=callgrind --trace-children=yes  
./bin/peloton &> /dev/null&
```

KCACHEGRIND

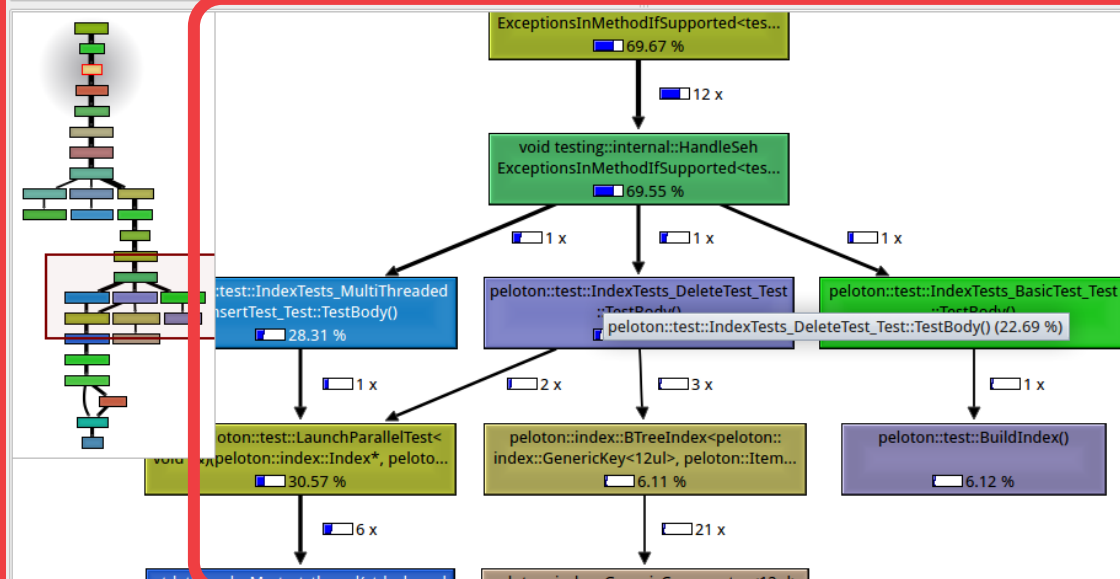
Profile data visualization tool

```
$ kcachegrind callgrind.out.12345
```

Cumulative Time Distribution

Incl.	Self	Called	Function
92.84	0.00	(0)	0x00000000000012d0
78.31	0.01	1	_dl_start
78.30	0.01	1	_dl_sysdep_start
78.29	0.02	1	dl_main
76.43	11.20	13	_dl_relocate_object
68.77	38.98	7 312	_dl_lookup_symbol_x
29.79	23.17	7 312	do_lookup_x
11.15	0.00	1	0x00000000000406f7e
11.14	0.00	1	(below main)
7.16	0.00	4	start_thread
7.04	0.00	6	0x0000000000008d370
7.03	0.00	6	std::thread::Impl<std::Bin...
6.95	0.52	319	peloton::index::GenericCom...
6.58	0.00	1	main
6.39	0.00	1	RUN_ALL_TESTS()
6.39	0.00	1	testing::UnitTests::Run()
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	bool testing::internal::Handl...
6.38	0.00	1	testing::internal::UnitTestIm...
5.73	3.57	7 149	check_match.9458
5.49	0.04	5	peloton::test::InsertTest(pel...
5.34	0.00	1	testing::TestCase::Run()
5.07	0.00	3	testing::TestInfo::Run()
4.93	0.04	46	peloton::index::BTreeIndex...
4.82	0.07	46	std::btree<peloton::index::G...
4.60	0.00	12	void testing::internal::Handl...
4.60	0.00	3	testing::Test::Run()
4.60	0.00	12	void testing::internal::Handl...
3.91	0.07	301	_dl_runtime_resolve
3.85	0.00	1	_libc_csu_init
3.85	0.27	301	_dl_fixup
3.69	1.06	2 187	malloc
3.59	0.53	2 024	peloton::Value::Value(pelot...
3.38	0.00	1	_dl_init
3.38	0.02	13	call_init.part.0

Callgraph View



CHOICE #2: PERF

Tool for using the performance counters subsystem in Linux.

- **-e** = sample the event cycles at the user level only
- **-c** = collect a sample every 2000 occurrences of event

```
$ perf record -e cycles:u -c 2000  
./tests/skiplist_index_test
```

Uses counters for tracking events

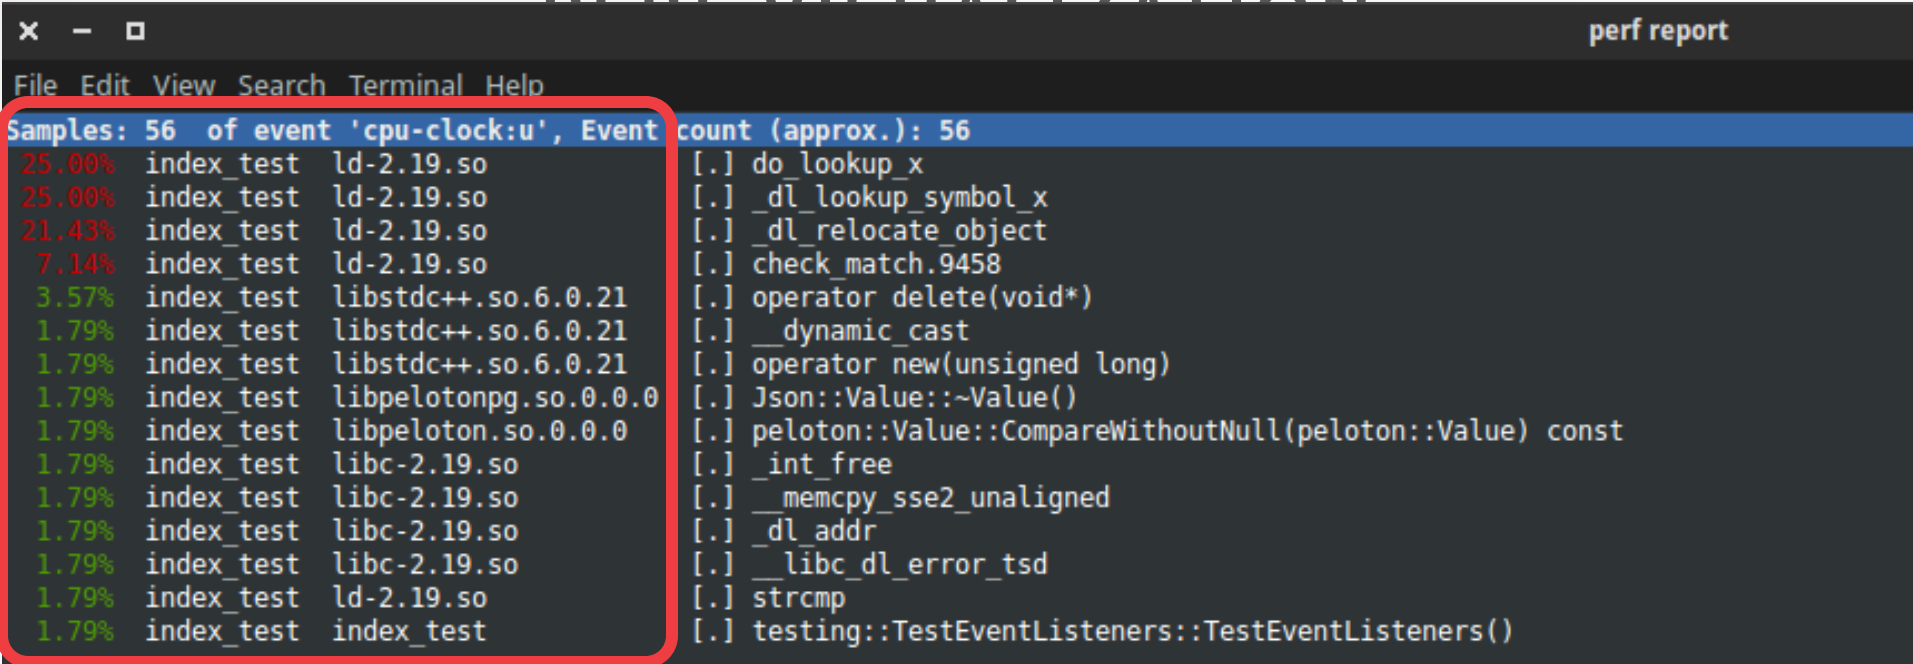
- On counter overflow, the kernel records a sample
- Sample contains info about program execution

PERF VISUALIZATION

We can also use **perf** to visualize the generated profile for our application.

```
$ perf report
```

PERF VISUALIZATION



Cumulative Time
Distribution

PERF EVENTS

Supports several other events like:

- L1-dcache-load-misses
- branch-misses

To see a list of events:

```
$ perf list
```

Another usage example:

```
$ perf record -e cycles,LLC-load-misses -c 2000  
./tests/skiplist_index_test
```


REFERENCES

Valgrind

- [The Valgrind Quick Start Guide](#)
- [Callgrind](#)
- [Kcachegrind](#)
- [Tips for the Profiling/Optimization process](#)

Perf

- [Perf Tutorial](#)
- [Perf Examples](#)
- [Perf Analysis Tools](#)

NEXT CLASS

Indexing for OLAP workloads.
→ More from Microsoft Research...